

DEEP DETERMINISTIC POLICY GRADIENT FOR CONTINUOUS CONTROL

HUBERT DE LASSUS

Training Multiple Agents To Solve A Continuous Control Task Using Deep Deterministic Policy Gradient.

For this toy example, the problem is to solve in UNITY ML environment the task of letting 20 robotic arms each keep control of its ball as it pushes the ball around. First, let's define some terminology. The world we address is defined as a Markovian Decision Process MDP with states s , action a , reward r and the conditional Probability P of future state s' given current state s , action a , reward r and discount factor γ ."

Environment: The conditional distribution of the state transitions and the reward function constitute the model of the environment.

Action: set A of available move from a given state s to a state s' . Only the elements of the set A where $P(s'|s, a) > 0$ are considered.

Episode: An episode is a complete sequence of events from an initial state to a final state.

Terminal states: The states that have no available actions are called terminal states.

Cumulative reward: The cumulative reward is the discounted sum of rewards accumulated throughout an episode.

Policy: A Policy is the agent's strategy to choose an action at each state.

Optimal policy: the policy that maximizes the expectation of cumulative reward.

The problem is modeled as an episodic task during which the agents have to maximize the expected cumulative rewards.

Because we choose to model the solution as a DDPG algorithm, let us introduce its concept.

DDPG algorithm concept. This is an algorithm that lies in between Value based methods and Policy based methods. While actor function specifies action a given the current state of the environment s , critic value function specifies a Temporal Difference (TD) Error to criticize the actions made by the actor.

- Stochastic methods are of the form: $\pi_\theta(a|s) = \mathbb{P}[a|s; \theta]$
- Deterministic methods are of the form: $a = \mu_\theta(s)$
- Computing stochastic gradient requires more samples, as it integrates over both state and action space. Deterministic gradient is preferable as it integrates over state space only.
- In DQN, action was selected as: $a_t = \max_a Q^*(\phi(s_t), a; \theta)$
- Above equation is not practical for continuous action space. Using deterministic policy allows us to use: $a_t = \mu(s_t|\theta_\mu)$. But, deterministic policy gradient might not explore the full state and action space. To overcome this, we introduce a noise process N :

$$a_t = \mu((s_t|\theta_\mu) + N_t)$$

This replaces the epsilon greedy algorithm of DQN for state and action space exploration. We want to maximize the rewards (Q-values) received over the sampled mini-batch. The gradient is given as:

$$\nabla_{\theta_\mu} J \approx \mathbb{E}_{s_t} \sim \rho^\beta [\nabla_{\theta^\mu} Q(s, a|\theta^Q)|_{s=s_t, a=\mu(s_t|\theta^\mu)}]$$

Applying the chain rule using μ instead of Q we get

$$\nabla_{\theta_\mu} J = \mathbb{E}_{s_t} \sim \rho^\beta [\nabla_a Q(s, a|\theta^Q)|_{s=s_t, a=\mu(s_t)} \nabla_{\theta_\mu} \mu(s|\theta^\mu)|_{s=s_t}]$$

This equation yields the maximum expected reward as we update the parameters using gradient ascent.

DDPG Algorithm steps. (This analysis is borrowed from San Jose University CS DDPG course hand out.)

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ

Initialize replay buffer R

Initialize target networks Q' and μ' with weights

$$\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$$

This step use target network to do off-policy updates

for episode = 1, M do

Initialize a random process \mathcal{N} for action exploration.

Receive initial observation state s_1

for $t = 1, T$ do

Select action using deterministic actor $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$

do it according to the current policy and exploration noise.

Execute action a_t and observe reward r_t , then observe new state s_{t+1} .

Store transition (s_t, a_t, r_t, s_{t+1}) in R .

Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R .

This is the experience replay sequence.

Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta_\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for

end for

Learning Algorithm. The actor implements a Policy gradient algorithm

The critic implements a Q-Learning algorithm with three main steps:

- 1) A sample step
- 2) A learn step
- 3) A model change step

To sample the environment a Multi Layer Perceptron is used to estimate the value actions based on environment observations.

Experience replay is used to reduce the oscillations of the output function of the network and accelerate the learning process by emphasizing the most meaningful samples. During the Learning step, a batch of past experiences is randomly sampled to train the agent. The randomness of this experiences selection, helps also the learning process by reducing correlation between input samples.

Hyperparameters. The parameters pertaining to the algorithms are

- Replay buffer size: int(1e6)
- Minibatch size: 64
- Discount factor: 0.99
- Learning rate actor : 1e-4
- Learning rate critic : 3e-4
- L2 Weight decay : 0.0001
- τ soft update of target parameters: 1e-3
- Noise level μ, θ, σ

The parameters pertaining to the actor network are

- Number of layers and nodes: 1 input layer with 33 nodes, 2 hidden layers with 64 fully connected nodes, 1 output fully connected layer layer with 4 nodes.

- Activation function: hyperbolic tangent of ReLu

- The neural network optimizer: Adam optimizer and its parameters.

The parameters pertaining to the critic network are

- Number of layers and nodes: 1 input layer with 33 nodes, 2 hidden layers with 64 fully connected nodes, 1 hidden layer with 32 fully connected nodes, 1 output fully connected layer layer with 1 node.

- Activation function: hyperbolic tangent of ReLu

- The neural network optimizer: Adam optimizer and its parameters.

Actor Model Architecture. The actor model architecture is defined by 5 variables: *StateSize* is the input dimensions of the network *ActionSize* is the output dimensions of the network Seed will initialize the weights of the network *fc1units* and *fc2units* are the number of nodes in the hidden layers of the network

The input layer has 33 nodes corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints – every action vector value is between $[-1, 1]$.

The first hidden layer has 64 and the second 64 nodes. The output layer has 4 nodes one for each action. Optimizing the number of nodes in the network leads to faster learning and better generalization. By setting the number of hidden layers to two and using dichotomy node sampling, I eventually found that 64 hidden nodes on two hidden layers give better results than a larger network.

Critic Model Architecture. The critic model architecture is defined by 5 variables: *StateSize* is the input dimensions of the network *ActionSize* is the output dimensions of the network Seed will initialize the weights of the network *fc1units*, *fc2unit* and *fc3units* are the number of nodes in the hidden layers of the network. the first two layers have 64 nodes, the third layer has 32 nodes. *fc4units* is the output layer with one output node.

The input layer has 33 nodes corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints – every action vector value is between $[-1, 1]$.

Optimizing the number of nodes in the network to a minimum leads to faster learning and better generalization. I eventually found that 64-64-32 hidden nodes on two hidden layers give similar or better results than a larger network.

Training. The training runs for 600 episodes with 500-time steps per episode. When the average of the 100 most recent scores reaches 30 the problem is considered solved and the algorithm stops.

A training episode is conducted as follows:

- 1) Get predicted next-state actions from actor target models

- 2) Get Q values from critic for this next-state
- 3) Compute Q targets for current states (y_i),

$$Q_{targets} = rewards + (\gamma * Q_{targetsnext} * (1 - done))$$

- 4) Compute critic loss
- 5) Minimize the critic loss by backpropagation with adam optimizer
- 6) Minimize the actor loss by backpropagation with adam optimizer of mean critic loss
- 7) Update the target networks policy and value parameters using given batch of experience tuples.

$$Q_{targets} = r + \gamma * critictarget(nextstate, actortarget(nextstate))$$

where: $actortarget(state) \rightarrow action$ $critictarget(state, action) \rightarrow Qvalue$

Scores Dynamic. This DDPG algorithm with 20 agents has a smooth, stable learning curve that increases exponentially.

The scores achieved during the learning process are:

Episode 100 Average Score: 5.84

Episode 200 Average Score: 27.47

Episode 300 Average Score: 67.35

Episode 400 Average Score 118.22

Episode 500 Average Score 172.33

Episode 600 Average Score 232.09

Environment solved in 209 episodes! Average Score: 30.31

End Score: 232.09

Ideas For Future Work. There are many different ways this project can be improved. However, the problem is already solved in a stable fashion and it is not warranted that resource hungry refinements are necessary in this case. However, it should be possible to train faster. This is the domain where experiments would be useful using flips of states and noise. Other types of algorithms could be tested, including Proximal Policy Optimization, Trust region Policy, and prioritized experience replay.

Plot of Rewards. Figure 1

Figure 1 shows the scores during training.




FIGURE 1. Scores.